

Semantic Secrets: Private Authentication from Reconstructable Visual Semantics

Anonymous Author(s)

Abstract

This planning manuscript defines a revised research direction; it does not report a completed authentication system. We ask whether a remembered visual concept can be reconstructed through independently generated images, converted locally into a canonical typed semantic graph, and privately verified under a system-derived policy. P8 narrowed the prospective contribution after finding prior art for the component mechanisms. P9-v2 then tested two frozen local extractors that each had to infer a scene and author the complete exact credential JSON. Each produced an invalid graph on its first formal validation fixture. With 32 validation fixtures, one invalid output caps best-case schema validity at $31/32 = 96.875\%$, below the preregistered 98% requirement. Gate V2-A failed; full semantic accuracy was not estimated. P9-v3A now preregisters a distinct modular architecture: perception models emit bounded observations and provenance, while a deterministic compiler emits a canonical graph or typed failure. It freezes two pipelines, controlled and naturalistic capability strata, per-type credential eligibility, exact compiler tests, and new extraction/reconstruction gates. No v3 model, image, inference, authentication, or cryptographic experiment has run. P10 remains blocked.

1 Problem and scope

Authentication based on approximate semantic meaning has an unusual tension. Benign reconstructions must vary, yet every tolerated variation enlarges the set an attacker may submit. If a server stores a deterministic semantic verifier, the same structure that permits fuzzy acceptance may reveal semantic content, support offline candidate testing, or link reuse across services. Secure computation may hide the comparison while leaving the accepted region easily enumerable.

We study a prospective client-side pipeline in which a person reconstructs a remembered visual concept using free natural language. A local generator produces a new image, and a local extractor maps that image into a canonical semantic representation. The strongest target sends neither prompt, image, plaintext graph, nor raw embedding to the authentication infrastructure.

No participant study is planned. Researcher-authored concepts and reconstructions can measure only technical semantic reconstruction stability. They cannot establish memorability, recall, usability, preference, or natural credential choice.

2 The prompt and image are not the secret

Let C denote a remembered visual concept and P a free-form reconstruction prompt. Different expressions may aim at the same concept—for example, “a cat beside a red car” and “a feline next to a crimson automobile.” The hypothesis is not exact prompt reproduction. With generator G and randomness r ,

$$I = G(P, r), \quad (1)$$

where I is a transient, independently generated image. Pixels do not contribute password entropy merely because generation is stochastic. A versioned modular pipeline produces observations

$$O = \text{Observe}(I), \quad (2)$$

and a deterministic semantic compiler produces

$$S = C(O) \quad (3)$$

or a typed failure. Models infer bounded evidence; code owns schema, identifiers, duplicates, inverse normalization, derived counts and geometry, ordering, and serialization. The credential source is the canonical semantics and, operationally, the complete set of candidate semantics the verifier accepts.

This framing differs from treating direct text embeddings or a single generated image as the authentication object. It also creates a falsifiable condition: if independent generation and fresh extraction do not preserve security-relevant structure, the visual reconstruction hypothesis fails.

3 Canonical policy-aware semantics

We provisionally represent S as a typed graph

$$S = (V, E_S, A, C_S), \quad (4)$$

where V contains entities, E_S typed actions/relations, A attributes, and C_S counts and scene/context facts. A deterministic, versioned policy derivation function produces

$$\Pi(S) = (M, T), \quad (5)$$

where M contains mandatory security anchors and T contains secondary facts with bounded tolerance. The primary non-human study will not ask users to nominate M or T . Selection must be frozen from independent semantic rationale, training-only statistics, extractor reliability, attacker-frequency evidence, and protocol compatibility.

A P8-frozen reference functionality has the form

$$\text{Accept}_{\text{ref}} = \text{MandatoryOK}(M, M') \wedge [\text{ToSim}(T, T') \geq \tau], \quad (6)$$

The v3 observation/compiler contract is specified in docs/formal_specification_v3.md. Its broad candidate language L_{visual} contains entities, attributes, actions/interactions, derived geometry/counts, and scenes. A type enters $L_{\text{cred}} \subseteq L_{\text{visual}}$ only after independent capability evidence passes in both controlled and naturalistic strata. Actions and interactions are independently gated: neither is required nor silently discarded. The M/T notation and policy split are preserved but are not implemented or tuned in P9-v3A and are not claimed as standalone novelty.

4 Acceptance-region security

For an enrolled state, define the accepted region

$$\mathcal{A}(S, M, T) = \{(S', M', T') : \text{Accept}_{\text{ref}} = 1\}. \quad (7)$$

An attacker succeeds by finding any member, not by recovering the enrolled prompt, image, or exact graph. We distinguish:

- K_0 : generic public knowledge and random candidates;
- K_1 : population or frequency knowledge without target facts;
- K_2 : partial target-specific semantic information;
- K_3 : strong near-secret information.

For strategy Q , budget B , knowledge K_i , and observable view V , the primary quantity is

$$\text{Succ}(Q, B, K_i, V) = \Pr[\exists j \leq B : q_j \in \mathcal{A}(S, M, T)]. \quad (8)$$

The evaluation should report success@1, @5, @10, and @B, guesses to success, duplicate policy, uncertainty, and cost. Random FAR is only one distributional statistic. Frequency, targeted neighbours, partial knowledge, LLM/VLM/generator assistance, and adaptive semantic search may produce different risk. AI-assisted password guessing is established prior art [21]; this project must contribute a semantic acceptance-region analysis rather than the general idea of using a model to guess.

5 Privacy target

The protected record is abstractly

$$R = \text{Protect}(S, M, T, \text{account}, \text{domain}, \text{version}). \quad (9)$$

A later protocol should reproduce the frozen plaintext decision while revealing no more than its declared output and metadata. Candidate goals are:

- template confidentiality against semantic recovery or accepted-equivalent reconstruction;
- resistance to cheap database-only candidate validation under an exact stolen-record view;
- policy privacy for anchor identities, operators, weights, thresholds, and graph shape;
- unlinkable, domain-separated records and transcripts;
- transcript minimisation, replay/session binding, and explicit adaptive-query leakage;
- containment under named database, server, key/share, collusion, and total-compromise states.

Established fuzzy extractors [4], fuzzy aPAKE [5], OPRFs [3], OPAQUE [2], exact PSI [6, 20], fuzzy PSI/private biometric matching [1, 24], committed PSI [7], private fuzzy-record computation [9], private graph operations [12], and homomorphic vector matching [15] are candidate foundations, not novel primitives. Raw learned representations are not automatically private; embedding leakage is an explicit concern [22]. Protected-template unlinkability and irreversibility must be attacked rather than inferred from an encoding [8].

Semantic correctness and cryptographic privacy are independent. FAR/FRR do not prove privacy; private computation does not repair an easily guessed acceptance region.

6 Closest work and VSA boundary

VSA [10] is the closest mandatory baseline. It enrolls a reference image, extracts objects, CLIP-derived attributes, quantities, visibility, and 2x2 quadrant location, and lets a user select two or three semantic facts plus quantity operators. Authentication uses an arbitrary satisfying image and an additional password. VSA serialises

canonical semantic tokens, evaluates Flexible Range Logic conjunctions, hashes most semantic fields together with the password using SHA-256, and retains quantity/operator metadata separately.

These features remove image-independent visual-semantic authentication, VLM extraction, canonical tokens, object, attribute, count, and quadrant rules, flexible quantity policies, conjunction, and user-selected semantic policy from our novelty space. We require two fair baselines: (i) VSA’s semantic-policy component for functional discrimination and (ii) the full semantic-plus-password hash architecture for security comparison. We must not compare a semantic-only system directly with VSA’s claimed combined semantic-plus-password space as though the credentials were equivalent.

VSA calls its hash binding server-opaque and explicitly weaker than strict zero knowledge. It acknowledges small semantic inputs and fast SHA-256 offline guessing, recommending Argon2id with a user-specific salt. Its COCO-based positive FRR samples are automatically selected from cached detections rather than independently re-inferred. It reports no human study and leaves Semantic Relationship Logic to future work. It does not target private/unlinkable verification.

SemanticLock uses semantically linked image tokens and story structure for graphical authentication [19]. PassStyles [17] and Omokage [18] use StyleGAN-generated changing face candidates and recognition criteria, preventing broad “first semantic”, “first generative”, or “first changing-image” authentication language. Adaptive score-guided verifier search is likewise established [16]. The P8 review found no verified source combining free-language independent text-to-image reconstruction and fresh typed extraction with a system-derived policy, finite-budget semantic attacks, and construction-specific private/domain-separated verification. Gate V2-N passed for that scoped question only; this is not an exhaustive absence or first claim.

The v3 shortlist is grounded in prior visual-understanding work. Grounding DINO produces open-set, language-conditioned boxes [14]; SigLIP 2 provides contrastive vision-language encoders for closed-label scoring and localization-oriented transfer [23]; and SGTR predicts entity/predicate proposals and graph connectivity end to end [13]. Visual Genome illustrates the dense object/attribute/relation annotation problem and its dataset limitations [11]. These are established perception components, not authentication or privacy contributions.

7 Research questions

RQ1—Reconstruction. Can independently generated images representing the same intended visual concept produce stable canonical semantic credentials?

RQ2—Policy-aware matching. Can a system-derived $S = (M, T)$ policy preserve legitimate technical reconstruction variation while rejecting security-critical semantic differences?

RQ3—Acceptance-region security. How vulnerable is the induced accepted region under K_0 – K_3 attackers and finite budgets?

RQ4—Template/privacy leakage. What semantics can be inferred from stored records, policies, and authentication transcripts under named compromise views?

RQ5—Private verification. Can the frozen policy be evaluated while hiding the semantic credential and resisting practical database-only validation and cross-domain linkage under stated assumptions?

RQ6—Practicality. What computation, bandwidth, latency, storage, trust, and deployment assumptions are required?

8 Proposed methodology and baselines

The execution order is deliberately staged.

Novelty and preregistration. P8 completed a focused primary-source novelty refresh and froze graph tasks, controlled concepts, splits, capability screening, policy derivation, attack budgets, uncertainty, and numeric gate thresholds. It performed no model acquisition, dataset acquisition, v2 inference, sealed-v1 access, or cryptographic implementation.

Representation before authentication. P9-v2 screened two pre-registered monolithic extractors on 96 controlled images. Both failed the schema-validity prerequisite under exact logical futility, and reconstruction was not executed. P9-v3A is a new architecture/preregistration phase, not a repair. It freezes $I = G(P, r)$, $O = \text{Observe}(I)$, $S = C(O)$ and $(M, T) = \Pi(S)$; two modular pipelines; a new 240-image capability design; compiler/perception metrics; and Gates V3-A1/V3-A2. It runs no model.

Policy and attackers. P10 compares dense similarity, global fuzzy overlap, a feasible VSA-style policy, and the proposed system-derived M/T policy. P11 evaluates the full acceptance region under budgeted attackers before any cryptographic implementation.

Private verification. Only a constructive semantic-security gate permits P12 to define the ideal functionality and compare mature constructions. Minimal POCs precede selection. The selected protocol is then evaluated for reference correctness, leakage, offline validation, compromise, linkage, and cost.

Baselines. Planned baselines include direct-text processing, dense visual/text similarity, global fuzzy semantic overlap, VSA semantic-policy behaviour, the full VSA binding model where reproducible, plaintext/hash records, and relevant mature private-matching constructions. Any inaccessible or underspecified VSA behaviour will be documented rather than guessed.

9 P9 capability outcome

The P9 capability manifest contains 64 development and 32 validation fixtures, equally divided between procedural/composite and researcher-authored synthetic-raster scenes. It accesses no participant or sealed-v1 data. Moondream2 at revision prefix 9a7d40240508 emitted no JSON on formal validation fixture cap-v2-064 in 157.57 seconds with 4.58 GiB peak process RSS. SmolVLM2-2.2B-Instruct at revision prefix c0a7af506d0f emitted malformed, truncated JSON in 392.10 seconds with 5.62 GiB peak RSS. The artifact records both complete immutable revisions.

The preregistration did not specify an early-stop rule. After separate plumbing smoke and before formal validation output, we implemented exact logical futility: because all checks are conjunctive, one invalid result makes the 98% schema-validity check impossible even if all remaining fixtures pass. This procedural deviation

cannot create a pass, but it limits the negative evidence. We did not estimate full atom F1, bootstrap intervals, determinism, structured error strata, or full-set latency medians. We made no candidate replacement, prompt/output repair, fine-tuning, or authentication-outcome selection.

10 P9-v3A modular extraction preregistration

The v2 failure conflated serialization correctness with semantic correctness. A model can describe an image correctly yet truncate JSON, or emit valid JSON containing false facts. V3 therefore makes the observation object explicit: detections carry category, normalized box, component-local confidence and revision; attributes, unary actions, binary interactions, and scene hypotheses reference those detections and retain provenance, abstention, failure, latency and memory events. Confidence values are not compared across components or treated as calibrated probabilities.

The deterministic compiler validates versions, labels, boxes and references; resolves same-category duplicates under a frozen IoU/tie-break rule; assigns canonical identifiers; derives counts and spatial relations; normalizes inverse edges; and serializes sorted UTF-8 JSON. It returns a valid graph or a valid typed failure, never malformed output. At least 320 exact tests must pass with 100% oracle agreement, schema validity and byte repeatability. One compiler invariant failure blocks V3-A1 regardless of perception scores.

The two frozen pipelines test distinct assumptions. The first combines Grounding DINO Tiny with SigLIP 2 Base for entity localization and closed-label crop, pair-crop and scene scoring. The second combines SGTR's graph-native entity/predicate assembly with SigLIP 2 attributes and scenes. Exact revisions and licence conditions are preregistered; missing SGTR checkpoint terms fail closed rather than opening a replacement slot. Planning bounds permit one GPU with at most 24 GiB VRAM and 32 GiB RAM, median latency at most 30 seconds/image and p95 at most 60 seconds/image. These are thresholds, not measurements.

Future P9-v3B uses 120 controlled geometric/composited fixtures and 120 new naturalistic outputs from the intended local generator. Each stratum has 24 semantic families split 12/12 between development and validation, five images per family. Labels and provenance are fixed before model output. An atom type must pass precision, recall, F1, coverage and family-level uncertainty separately in both strata; pooled results cannot rescue failure. Gate V3-A1 requires a complete pipeline with eligible entity plus at least two further types, including one structural count, geometry, action, or interaction type. Types from different pipelines cannot be unioned.

P9-v3C is future and requires a separate reconstruction freeze. Gate V3-A2 requires eligible enrolment, FRR, targeted-neighbour FAR, random FAR, same/near AUC, and a material paired M/T improvement over the strongest baseline. Passing it would establish only technical reconstruction feasibility. P10 remains blocked until V3-A2; cryptography remains further downstream.

11 Decision gates

Gate V2-N passed with narrowing. Gate V2-A is the immutable failed monolithic path. The explicit new-version path uses V3-A1 modular extraction viability and V3-A2 independent reconstruction

viability before the retained V2-B policy, V2-C acceptance-region, V2-D private-functionality, V2-E protocol/privacy, and V2-F end-to-end gates. No later cost may bypass an earlier failure. Valid outcomes include standalone authentication, second factor, policy-constrained credential, measurement/negative contribution, or stop.

12 Formative v1 evidence

The completed P0–P7 programme evaluated a different system, visual-semantic-pipeline-v1. P6’s training-selected weighted matcher had validation FRR 20.8%, targeted-neighbour FAR 75.0%, and random FAR 3.0%; all tested representations failed the original gate. Twelve v1 test families remain sealed. P7 then found no uncertainty-supported material benefit from the tested SD-Turbo/Florence/SigLIP image pathways and removed that image stage from the v1 authentication core.

These are immutable formative negative results, not v2 measurements. They motivate explicit security anchors and independent capability screening but may not tune the v2 policy, representation, data, or gates. The v2 hypothesis requires new identifiers, preregistration, and evidence.

13 Expected contributions and limitations

If all gates pass, C1 would establish technical behaviour of the scoped free-language generative reconstruction path; C2 would establish only the measured incremental value of the frozen system-derived typed policy; C3 would characterise the accepted region under finite-budget attackers; and C4 would establish construction-specific private and unlinkable verification. C3/C4 carry the prospective principal contribution. P8 rejected C5 as standalone and retained adaptive semantic verification-oracle leakage as mandatory evidence inside C3/C4.

The frozen v2 path stopped at the first risk: neither tested extractor met the structured-output prerequisite. P9-v3A provides an independently justified representation architecture but no empirical rescue. Whether either modular pipeline works remains unresolved, as do reconstruction, policy, guessing-region, and private-construction questions. The controlled methodology cannot establish human memorability or usability. A future participant study would require separate ethics, design, and evidence.

14 Conclusion

This work treats visual generation as a testable reconstruction layer, not a source of magical entropy, and privacy as a construction-specific property rather than a consequence of hashing semantics. P9-v2 falsified its frozen monolithic extractor path, and that evidence remains unchanged. P9-v3A supplies the required explicit reframe: models produce evidence, a deterministic compiler constructs the credential, and atom types qualify independently on new controlled and naturalistic data. It passes no empirical gate. P9-v3B is the next permitted experiment only after explicit authorisation; P9-v3C, P10, attacker evaluation, and cryptographic prototyping remain blocked by their gates.

References

- [1] Erik-Oliver Blass and Guevara Noubir. 2026. Assumption-Free Fuzzy PSI via Predicate Encryption. In *35th USENIX Security Symposium (USENIX Security*

- 26). USENIX Association, Baltimore, MD. <https://www.usenix.org/conference/usenixsecurity26/presentation/blasse>
- [2] D. Bourdreux, H. Krawczyk, K. Lewi, and C. A. Wood. 2025. *The OPAQUE Augmented Password-Authenticated Key Exchange (aPAKE) Protocol*. RFC 9807. RFC Editor. <https://www.rfc-editor.org/rfc/rfc9807.html> Informational RFC; CFRG/IRTF consensus document.
- [3] A. Davidson, A. Faz-Hernandez, N. Sullivan, and C. A. Wood. 2023. *Oblivious Pseudorandom Functions (OPRFs) Using Prime-Order Groups*. RFC 9497. RFC Editor. <https://www.rfc-editor.org/rfc/rfc9497.html> Informational RFC; CFRG/IRTF consensus document.
- [4] Yevgeniy Dodis, Rafail Ostrovsky, Leonid Reyzin, and Adam Smith. 2008. Fuzzy Extractors: How to Generate Strong Keys from Biometrics and Other Noisy Data. *SIAM J. Comput.* 38, 1 (2008), 97–139. <https://www.cs.bu.edu/~reyzin/fuzzy.html>
- [5] Andreas Erwig, Julia Hesse, Maximilian Ortl, and Siavash Riahi. 2020. Fuzzy Asymmetric Password-Authenticated Key Exchange. In *Advances in Cryptology – ASIACRYPT 2020*. Springer. <https://research.ibm.com/publications/fuzzy-asymmetric-password-authenticated-key-exchange>
- [6] Michael J. Freedman, Kobbi Nissim, and Benny Pinkas. 2004. Efficient Private Matching and Set Intersection. In *Advances in Cryptology – EUROCRYPT 2004 (Lecture Notes in Computer Science, Vol. 3027)*. Springer, 1–19. <https://crypto.stanford.edu/portia/pubs/articles/FNP227785321.html>
- [7] Aarushi Goel, Peihan Miao, Phuoc Van Long Pham, and Satvinder Singh. 2026. PICS: Private Intersection over Committed (and reusable) Sets. In *35th USENIX Security Symposium (USENIX Security 26)*. USENIX Association, Baltimore, MD. <https://www.usenix.org/conference/usenixsecurity26/presentation/goel>
- [8] Marta Gomez-Barrero, Christian Rathgeb, Javier Galbally, Christoph Busch, and Julian Fierrez. 2016. Unlinkable and Irreversible Biometric Template Protection Based on Bloom Filters. *Information Sciences* 370–371 (2016), 18–32. doi:10.1016/j.ins.2016.06.046
- [9] KyooHyung Han, Seongkwang Kim, and Yongha Son. 2025. Private Computation on Common Fuzzy Records. *Proceedings on Privacy Enhancing Technologies* 2025, 1 (2025), 567–583. doi:10.56553/popets-2025-0031
- [10] Pil-seong Jeong. 2026. Design and Security Validation of Image-Agnostic Visual Semantic Authentication (VSA) Framework Based on Vision-Language Model. *Journal of The Korea Institute of Information Security and Cryptology* 36, 3 (June 2026), 935–948. doi:10.13089/JKIISC.2026.36.3.935 Peer-reviewed article in Korean.
- [11] Ranjay Krishna, Yuke Zhu, Oliver Groth, Justin Johnson, Kenji Hata, Joshua Kravitz, Stephanie Chen, Yannis Kalantidis, Li-Jia Li, David A. Shamma, Michael S. Bernstein, and Li Fei-Fei. 2017. Visual Genome: Connecting Language and Vision Using Crowdsourced Dense Image Annotations. *International Journal of Computer Vision* 123 (2017), 32–73. doi:10.1007/s11263-016-0981-7
- [12] Peng Li, Fucai Zhou, Zifeng Xu, and Yuxi Li. 2020. Privacy-Preserving Graph Operations for Mobile Authentication. *Wireless Communications and Mobile Computing* (2020). doi:10.1155/2020/8859213
- [13] Rongjie Li, Songyang Zhang, and Xuming He. 2022. SGTR: End-to-End Scene Graph Generation with Transformer. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. IEEE, 19486–19496.
- [14] Shilong Liu, Zhaoyang Zeng, Tianhe Ren, Feng Li, Hao Zhang, Jie Yang, Qing Jiang, Chunyuan Li, Jianwei Yang, Hang Su, Jun Zhu, and Lei Zhang. 2024. Grounding DINO: Marrying DINO with Grounded Pre-Training for Open-Set Object Detection. In *Computer Vision – ECCV 2024*. Springer. <https://arxiv.org/abs/2303.05499>
- [15] Sam Martin, Nirajan Koirala, Helena Berens, Thomas Rozgonyi, Micah Brody, and Taehong Jung. 2025. HyDia: FHE-based Facial Matching with Hybrid Approximations and Diagonalization. *Proceedings on Privacy Enhancing Technologies* 2025, 4 (2025), 585–604. doi:10.56553/popets-2025-0146
- [16] Marcos Martinez-Diaz, Julian Fierrez, Javier Galbally, and Javier Ortega-Garcia. 2011. An Evaluation of Indirect Attacks and Countermeasures in Fingerprint Verification Systems. *Pattern Recognition Letters* 32, 12 (2011), 1643–1651. doi:10.1016/j.patrec.2011.04.005
- [17] Hikaru Matsuzaki, Tomoyuki Maekawa, Michita Imai, and Kentaro Ishii. 2025. PassStyles: Graphical Authentication Using Mixed Face Styles. *IEEE Access* 13 (2025), 165411–165428. doi:10.1109/ACCESS.2025.3612089
- [18] Hikaru Matsuzaki, Tomoyuki Maekawa, Michita Imai, and Kentaro Ishii. 2026. Omokage Authentication: A Graphical Authentication System Using Mixed Visual Resemblance. *IEEE Access* 14 (2026), 93009–93026. doi:10.1109/ACCESS.2026.3704981
- [19] Ilesanmi Olade, Hai-Ning Liang, and Charles Fleming. 2023. Story-Based Authentication for Mobile Devices Using Semantically-Linked Images. *International Journal of Human-Computer Studies* 171 (March 2023), 102967. doi:10.1016/j.ijhcs.2022.102967
- [20] Benny Pinkas, Thomas Schneider, Gil Segev, and Michael Zohner. 2015. Phasing: Private Set Intersection Using Permutation-Based Hashing. In *24th USENIX Security Symposium (USENIX Security 15)*. USENIX Association, Washington, DC, 515–530. <https://www.usenix.org/conference/usenixsecurity15/technical-sessions/presentation/pinkas>

- [21] Javier Rando, Fernando Perez-Cruz, and Briland Hitaj. 2023. PassGPT: Password Modeling and (Guided) Generation with Large Language Models. In *Computer Security – ESORICS 2023 (Lecture Notes in Computer Science, Vol. 14347)*. Springer, 164–183. doi:10.1007/978-3-031-51482-1_9
- [22] Congzheng Song and Ananth Raghunathan. 2020. Information Leakage in Embedding Models. In *Proceedings of the 2020 ACM SIGSAC Conference on Computer and Communications Security*. Association for Computing Machinery, 377–390. doi:10.1145/3372297.3417270
- [23] Michael Tschannen, Alexey Gritsenko, Xiao Wang, Muhammad Ferjad Naeem, Ibrahim Alabdulmohsin, Nikhil Parthasarathy, Talfan Evans, Lucas Beyer, Ye Xia, Basil Mustafa, Olivier Hénaff, Jeremiah Harmsen, Andreas Steiner, and Xiaohua Zhai. 2025. SigLIP 2: Multilingual Vision-Language Encoders with Improved Semantic Understanding, Localization, and Dense Features. arXiv:2502.14786 [cs.CV] <https://arxiv.org/abs/2502.14786> Preprint.
- [24] Erkam Uzun, Simon P. Chung, Vladimir Kolesnikov, Alexandra Boldyreva, and Wenke Lee. 2021. Fuzzy Labeled Private Set Intersection with Applications to Private Real-Time Biometric Search. In *30th USENIX Security Symposium (USENIX Security 21)*. USENIX Association, 911–928. <https://www.usenix.org/conference/usenixsecurity21/presentation/uzun>